Python Task

Manasa Barik

Prompt 1: Python Data Types

Act as a Python instructor.

Explain the basic Python data types: int, float, string, and boolean. Provide simple, well-commented Python examples for each. Explain where these data types are commonly used in real applications. Keep the explanation beginner-friendly and clear.

↳ 9 cells hidden

Prompt 2: Lists in Python

Act as a Python programmer.

Explain Python lists and their properties (ordered and mutable). Demonstrate list creation, indexing, updating values, appending elements, and iteration. Provide clean Python code with comments and simple explanations.

↳ 12 cells hidden

Prompt 3: Tuples in Python

Act as a Python developer.

Explain tuples in Python and how they differ from lists. Demonstrate tuple creation, indexing, and immutability with examples. Explain when tuples are preferred over lists.

Python tuples are ordered, immutable sequences of items. Like lists, they can store elements of different data types, but once a tuple is created, its elements cannot be changed, added, or removed. Tuples are defined by enclosing elements in parentheses `()`.

Key Differences from Lists:

Feature	Lists (`[]`)	Tuples (`()`)
Mutablity	Mutable (can be changed after creation)	Immutable (cannot be changed after creation)
Syntax	Square brackets `[]`	Parentheses `()`
Performance	Slower for iteration and access	Faster for iteration and access
Use Case	Collections of items that might change	Collections of related items that shouldn't change, c
Methods	Many methods (append, extend, insert, pop)	Fewer methods (count, index)

Let's dive into some examples.

1. Tuple Creation

Tuples are created by placing all the items (elements) inside parentheses `()`, separated by commas. A tuple with a single item needs a trailing comma.

```
# Creating an empty tuple
empty_tuple = ()
print(f"Empty tuple: {empty_tuple}, Type: {type(empty_tuple)}")

# Creating a tuple with mixed data types
my_tuple = (1, "banana", 2.718, False)
print(f"My tuple: {my_tuple}, Type: {type(my_tuple)}")

# Creating a tuple without parentheses (tuple packing) - Python infers it's
animal = "dog", "mammal", 4
print(f"Animal tuple (packed): {animal}, Type: {type(animal)}")

# Creating a single-element tuple (note the comma!)
single_element_tuple = (5,)
print(f"Single element tuple: {single_element_tuple}, Type: {type(single_ele

# A common mistake: without the comma, it's just an int in parentheses
not_a_tuple = (5)
print(f"Not a tuple: {not_a_tuple}, Type: {type(not_a_tuple)}")
```

```
Empty tuple: (), Type: <class 'tuple'>
My tuple: (1, 'banana', 2.718, False), Type: <class 'tuple'>
Animal tuple (packed): ('dog', 'mammal', 4), Type: <class 'tuple'>
Single element tuple: (5,), Type: <class 'tuple'>
Not a tuple: 5, Type: <class 'int'>
```

2. Indexing Tuples

Just like lists and strings, elements in a tuple can be accessed using indexing. The index starts from `0` for the first element. Negative indexing also works (`-1` for the last element).

```
# Using the 'my_tuple' from before: (1, 'banana', 2.718, False)
```

```
print(f"First element of my_tuple: {my_tuple[0]}")        # Output: 1
print(f"Third element of my_tuple: {my_tuple[2]}")         # Output: 2.718
print(f"Last element of my_tuple: {my_tuple[-1]}")         # Output: False

# You can also slice tuples to get a sub-tuple
sub_tuple = my_tuple[1:3] # Elements from index 1 up to (but not including)
print(f"Sliced tuple (my_tuple[1:3]): {sub_tuple}") # Output: ('banana', 2.7
```

```
First element of my_tuple: 1
Third element of my_tuple: 2.718
Last element of my_tuple: False
Sliced tuple (my_tuple[1:3]): ('banana', 2.718)
```

3. Immutability of Tuples

This is the most crucial property of tuples. Once created, you cannot change, add, or remove elements from a tuple. If you try, Python will raise a `TypeError`.

```
coordinates = (10, 20)
print(f"Original coordinates tuple: {coordinates}")

# Attempting to change an element will result in an error
try:
    coordinates[0] = 15
except TypeError as e:
    print(f"Error trying to modify a tuple element: {e}")

# Attempting to add an element will also result in an error
try:
    coordinates.append(30)
except AttributeError as e:
    print(f"Error trying to append to a tuple: {e}")

# If you need to 'change' a tuple, you actually create a new one
new_coordinates = (coordinates[0], 25) # Recreating with a 'changed' second
print(f"New coordinates tuple: {new_coordinates}")
```

```
Original coordinates tuple: (10, 20)
Error trying to modify a tuple element: 'tuple' object does not support item
Error trying to append to a tuple: 'tuple' object has no attribute 'append'
New coordinates tuple: (10, 25)
```

When are Tuples Preferred Over Lists?

Tuples are excellent choices in specific scenarios due to their immutability:

1. **Data that shouldn't change**: When you have a collection of items that should remain constant throughout the program's execution, like `(latitude, longitude)`, `(r, g, b)` color codes, or `(year, month, day)` for a fixed date. Their immutability acts as a guarantee against accidental modification.

2. **Function Arguments and Return Values**: Functions often accept or return multiple values. Tuples are a natural fit for this, especially when the order and number of elements are fixed (e.g., a function returning `(min_value, max_value)`).
3. **Dictionary Keys**: Because tuples are immutable, they can be used as keys in dictionaries, whereas lists cannot. This is useful for composite keys, such as `('Alice', 'Math')`.
4. **Performance Optimization**: Tuples are generally slightly faster and consume less memory than lists for similar data, especially for larger collections, because their size is fixed. While often a minor difference, it can be relevant in performance-critical applications.
5. **Data Integrity and Security**: In situations where data integrity is paramount, using tuples ensures that the data structure itself cannot be altered once defined, making the code more robust and predictable.

In essence, choose a tuple when the sequence of items is meant to be a fixed record, and its contents should not change. Choose a list when you need a collection that you expect to modify frequently.

Prompt 4: Dictionaries in Python

Act as a Python instructor.

Explain Python dictionaries and key-value pairs. Demonstrate dictionary creation, accessing values, updating entries, and iteration. Explain why dictionaries are useful for structured data. Provide simple, well-commented code examples.

↳ 12 cells hidden

Prompt 5: Object-Oriented Programming (OOP) in Python

Act as a Python software engineer.

Explain core OOP concepts in Python: class, object, attributes, and methods. Create a simple class and demonstrate object creation and method usage. Explain how OOP improves code organization and scalability. Use clear explanations and beginner-friendly examples.

↳ 6 cells hidden

Prompt 6: NumPy for Numerical Analysis

Act as a data scientist.

Explain what NumPy is and why it is used in Python. Demonstrate creating NumPy arrays and performing basic operations such as mean, sum, and max. Provide simple, well-commented examples and explain the output.

↳ 6 cells hidden

Prompt 7: Pandas for Data Analysis

Act as a data analyst.

Explain the purpose of Pandas in data analysis. Create a small sample DataFrame and demonstrate basic operations such as viewing data, filtering rows, and selecting columns. Explain how Pandas helps in analyzing structured data.

↳ 10 cells hidden

Prompt 8: Matplotlib for Data Visualization

Act as a Python data visualization expert.

Explain how Matplotlib is used for data visualization. Create a simple plot (bar or line chart) using sample data. Explain plot components such as axes, labels, and title. Describe why visualization is important in data analysis.

↳ 5 cells hidden

Prompt 9: Case Study using Synthetic Data

Act as a Python data analyst.

Create a small real-world case study where actual data is unavailable. Explain why synthetic (fake) data is commonly used during development and learning.

Using Python:

- Generate realistic synthetic data using a suitable library (NumPy random or Faker)
- Store the data in a Pandas DataFrame
- Perform basic analysis using Python and Pandas

- Visualize key insights using Matplotlib

Ensure the case study demonstrates how Python concepts such as data types, lists, dictionaries, NumPy, Pandas, and visualization work together in a practical scenario.

Keep the explanation beginner-friendly and the code well-commented.

Task

Here's the next step in our Python journey: a practical case study demonstrating how data types, lists, dictionaries, NumPy, Pandas, and Matplotlib come together to solve a real-world problem using synthetic data.

Why Synthetic Data?

In many development and learning scenarios, obtaining real-world data can be challenging due or impossible. Real data might be:

- **Proprietary or Confidential**: Companies cannot share sensitive customer, financial, or health information.
- **Unavailable**: For new products or systems, no historical data exists yet.
- **Too Complex/Messy**: Real data often requires extensive cleaning, which can distract from the core learning objectives of a new concept.
- **Limited**: Small datasets might not be sufficient to demonstrate the power of analytical tools.

Synthetic data is artificially generated data that mimics the statistical properties of real data. It allows developers and analysts to:

- **Test and Develop**: Build and test systems, algorithms, and analyses without legal or ethical concerns.
- **Learn and Experiment**: Practice data analysis, manipulation, and visualization techniques on realistic but safe datasets.
- **Demonstrate Concepts**: Illustrate how different tools and concepts work together, as we will do in this case study.
- **Overcome Data Scarcity**: Create large datasets for stress testing or training machine learning models when real data is scarce.

For our case study, we'll imagine we're an e-commerce analyst who wants to understand customer purchasing behavior but doesn't have access to actual customer transaction logs. We'll generate synthetic customer purchase data to simulate this scenario.

Case Study: Analyzing E-commerce Purchase Data with Python

We'll generate synthetic customer purchase data, store it in a Pandas DataFrame, perform some basic analysis, and visualize key insights.

Objective: To understand typical purchase patterns, popular product categories, and regional sales distribution.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from faker import Faker
import random
from datetime import datetime, timedelta

# Initialize Faker for generating realistic names and addresses
fake = Faker()

# --- 1. Generate Synthetic Data ---
num_transactions = 1000

# Lists of possible product categories and regions (implicitly using Python list
product_categories = ['Electronics', 'Books', 'Home & Kitchen', 'Apparel', 'Spor
regions = ['North', 'South', 'East', 'West', 'Central']

# Generate data for each transaction
data = []
for _ in range(num_transactions):
    customer_name = fake.name()

    # Use np.random for numerical data like purchase amount and quantity
    purchase_amount = round(np.random.uniform(5.0, 500.0), 2) # Float data type
    quantity = np.random.randint(1, 6) # Integer data type

    # Choose randomly from predefined lists (list data type)
    category = random.choice(product_categories)
    region = random.choice(regions)

    # Generate a random purchase date within the last year
    purchase_date = fake.date_time_between(start_date='-1y', end_date='now') # D

    # Store transaction as a dictionary (dictionary data type)
    data.append({
        'customer_name': customer_name,
        'product_category': category,
        'purchase_amount': purchase_amount,
        'quantity': quantity,
        'region': region,
```

```
        'purchase_date': purchase_date
    })

# --- 2. Create Pandas DataFrame ---
# The list of dictionaries is perfectly suited for creating a DataFrame
df_purchases = pd.DataFrame(data)

print("--- Sample of Synthetic Purchase Data ---")
display(df_purchases.head())
print(f"\nTotal transactions generated: {len(df_purchases)}")
print(f"Data types in DataFrame:\n{df_purchases.dtypes}")


# --- 3. Perform Basic Data Analysis ---

print("\n--- Basic Data Analysis ---")

# Calculate descriptive statistics for numerical columns
print("\nDescriptive Statistics for Purchase Amount and Quantity:")
display(df_purchases[['purchase_amount', 'quantity']].describe())

# Total sales amount
total_sales = df_purchases['purchase_amount'].sum()
print(f"\nTotal Sales Amount: ${total_sales:,.2f}")

# Group by product category to find total sales per category
sales_by_category = df_purchases.groupby('product_category')['purchase_amount'].
print("\nTotal Sales by Product Category:")
display(sales_by_category)

# Group by region to find average purchase amount per region
average_purchase_by_region = df_purchases.groupby('region')['purchase_amount'].m
print("\nAverage Purchase Amount by Region:")
display(average_purchase_by_region)

# Filter data: find all purchases above a certain amount (e.g., VIP purchases)
vip_purchases = df_purchases[df_purchases['purchase_amount'] > 400]
print(f"\nNumber of VIP Purchases (over $400): {len(vip_purchases)}")
if not vip_purchases.empty:
    print("Sample VIP Purchases:")
    display(vip_purchases.head())

# --- 4. Visualize Key Insights ---

print("\n--- Visualizing Key Insights ---")
```

```
# Visualization 1: Total Sales by Product Category (Bar Chart)
plt.figure(figsize=(12, 6))
sales_by_category.plot(kind='bar', color='lightcoral')
plt.title('Total Sales by Product Category')
plt.xlabel('Product Category')
plt.ylabel('Total Sales Amount ($)')
plt.xticks(rotation=45, ha='right')
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()

# Visualization 2: Distribution of Purchase Amounts (Histogram)
plt.figure(figsize=(10, 6))
plt.hist(df_purchases['purchase_amount'], bins=30, color='lightgreen', edgecolor
plt.title('Distribution of Purchase Amounts')
plt.xlabel('Purchase Amount ($)')
plt.ylabel('Number of Transactions')
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()

# Visualization 3: Average Purchase Amount by Region (Bar Chart)
plt.figure(figsize=(10, 6))
average_purchase_by_region.plot(kind='bar', color='skyblue')
plt.title('Average Purchase Amount by Region')
plt.xlabel('Region')
plt.ylabel('Average Purchase Amount ($)')
plt.xticks(rotation=0)
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()
```

Explain Synthetic Data

Subtask:

Provide an explanation of what synthetic data is and why it's commonly used in development and learning scenarios when real data is unavailable or sensitive.

Explain Synthetic Data

What is Synthetic Data?

Synthetic data is artificial data that is generated programmatically rather than collected from real-world events or measurements. It mimics the statistical properties and patterns of real data but does not contain any actual information from real individuals or systems. Essentially, it's 'fake' data that looks and behaves like real data.

Why is Synthetic Data Commonly Used?

Synthetic data is invaluable in development and learning scenarios, particularly when real data is unavailable, proprietary, or sensitive. Here are the key reasons for its widespread use:

1. **Privacy and Security**: Real-world data often contains sensitive information (e.g., personal identifiable information, financial records, health data). Using synthetic data allows developers and analysts to build, test, and demonstrate systems without compromising user privacy or violating data protection regulations (like GDPR or HIPAA).
2. **Unavailability or Scarcity of Real Data**: In many cases, especially during the early stages of a project or for rare events, real data might not yet exist or be scarce. Synthetic data can fill this gap, enabling progress even when real data collection is ongoing or impossible.
3. **Cost and Complexity of Real Data Collection**: Collecting, cleaning, and labeling real-world data can be incredibly time-consuming and expensive. Synthetic data generation can often be more efficient and cost-effective, particularly for large datasets.
4. **Development and Testing**: Developers use synthetic data to thoroughly test software applications, algorithms, and models across various scenarios, including edge cases that might be rare in real data. This ensures robustness and helps identify bugs early.
5. **Learning and Training**: For educational purposes or for training new data scientists and engineers, synthetic datasets provide a safe, reproducible, and easily shareable environment to practice data analysis, machine learning, and visualization techniques without dealing with the complexities and restrictions of real data.
6. **Reproducibility**: Experiments conducted with synthetic data are often easier to reproduce, as the data generation process can be precisely controlled and documented.
7. **Bias Mitigation**: In some cases, real datasets can contain biases. While synthetic data can also inherit these biases if generated poorly, it offers the

potential to create more balanced datasets to address and mitigate existing biases in models.

By using synthetic data, we can focus on the logic, functionality, and performance of our systems and analyses, making the development and learning process more efficient, ethical, and accessible.

Generate Synthetic Data

Subtask:

Generate realistic synthetic customer purchase data using libraries like NumPy for numerical data and Faker for textual data (e.g., customer names, product categories, purchase amounts, regions, and dates). The data types, lists, and dictionaries will be implicitly used during this step to structure the data.

Reasoning: I will import the necessary libraries, initialize Faker, define transaction parameters, and then loop to generate synthetic customer purchase data, storing each transaction as a dictionary in a list, as per the instructions.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from faker import Faker
import random
from datetime import datetime, timedelta

# Initialize Faker for generating fake data
fake = Faker()

# 3. Define the number of transactions to generate
num_transactions = 1000

# 4. Create Python lists for product_categories and regions
product_categories = ['Electronics', 'Clothing', 'Home Goods', 'Books', 'Foo
regions = ['North', 'South', 'East', 'West', 'Central']

# 5. Initialize an empty Python list called data to store transaction dictio
data = []

# 6. Loop num_transactions times to generate data
for _ in range(num_transactions):
    # a. Generate a customer_name
    customer_name = fake.name()

    # b. Generate a purchase_amount (float between 5.0 and 500.0) and round
    purchase_amount = round(np.random.uniform(5.0, 500.0), 2)

    # c. Generate a quantity (integer between 1 and 5)
```

```
        quantity = np.random.randint(1, 6) # randint is exclusive of the high en

        # d. Randomly select a category from product_categories
        category = random.choice(product_categories)

        # e. Randomly select a region from regions
        region = random.choice(regions)

        # f. Generate a purchase_date within the last year
        purchase_date = fake.date_time_between(start_date='-1y', end_date='now')

        # g. Create a dictionary for the current transaction and append it to th
        transaction = {
            'customer_name': customer_name,
            'product_category': category,
            'purchase_amount': purchase_amount,
            'quantity': quantity,
            'region': region,
            'purchase_date': purchase_date
        }
        data.append(transaction)

    print(f"Generated {num_transactions} synthetic transactions.")
    print("First 5 generated transactions:")
    for i in range(5):
        print(data[i])
```

```
Generated 1000 synthetic transactions.
First 5 generated transactions:
{'customer_name': 'Natalie Carroll', 'product_category': 'Home Goods', 'purch
{'customer_name': 'Margaret Wallace', 'product_category': 'Clothing', 'purcha
{'customer_name': 'Henry Morrison MD', 'product_category': 'Food', 'purchase_
{'customer_name': 'Theodore Cabrera', 'product_category': 'Food', 'purchase_a
{'customer_name': 'Denise Calderon', 'product_category': 'Electronics', 'purc
```

Reasoning: The previous cell failed because the `faker` library was not installed. I need to install it before it can be imported and used.

```
!pip install faker
print("Faker library installed.")
```

```
Collecting faker
  Downloading faker-40.4.0-py3-none-any.whl.metadata (16 kB)
Downloading faker-40.4.0-py3-none-any.whl (2.0 MB)
   ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 2.0/2.0 MB 67.1 MB/s eta 0:00:00
Installing collected packages: faker
Successfully installed faker-40.4.0
Faker library installed.
```

Reasoning: Since the `faker` library has been successfully installed, I will re-execute the code block to generate the synthetic customer purchase data as instructed.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from faker import Faker
import random
from datetime import datetime, timedelta

# Initialize Faker for generating fake data
fake = Faker()

# 3. Define the number of transactions to generate
num_transactions = 1000

# 4. Create Python lists for product_categories and regions
product_categories = ['Electronics', 'Clothing', 'Home Goods', 'Books', 'Foo
regions = ['North', 'South', 'East', 'West', 'Central']

# 5. Initialize an empty Python list called data to store transaction dictio
data = []

# 6. Loop num_transactions times to generate data
for _ in range(num_transactions):
    # a. Generate a customer_name
    customer_name = fake.name()

    # b. Generate a purchase_amount (float between 5.0 and 500.0) and round
    purchase_amount = round(np.random.uniform(5.0, 500.0), 2)

    # c. Generate a quantity (integer between 1 and 5)
    quantity = np.random.randint(1, 6) # randint is exclusive of the high en

    # d. Randomly select a category from product_categories
    category = random.choice(product_categories)

    # e. Randomly select a region from regions
    region = random.choice(regions)

    # f. Generate a purchase_date within the last year
    purchase_date = fake.date_time_between(start_date='-1y', end_date='now')

    # g. Create a dictionary for the current transaction and append it to th
    transaction = {
        'customer_name': customer_name,
        'product_category': category,
        'purchase_amount': purchase_amount,
        'quantity': quantity,
        'region': region,
        'purchase_date': purchase_date
    }
    data.append(transaction)

print(f"Generated {num_transactions} synthetic transactions.")
print("First 5 generated transactions:")
for i in range(5):
    print(data[i])
```

```
Generated 1000 synthetic transactions.
First 5 generated transactions:
{'customer_name': 'Candace Price', 'product_category': 'Food', 'purchase_amou
{'customer_name': 'Anthony Russell', 'product_category': 'Food', 'purchase_am
{'customer_name': 'Charles Nelson', 'product_category': 'Home Goods', 'purcha
{'customer_name': 'Christopher Zamora', 'product_category': 'Home Goods', 'pu
{'customer_name': 'Amanda Coleman', 'product_category': 'Home Goods', 'purcha
```

Reasoning: The synthetic data has been generated successfully as a list of dictionaries. The next logical step is to convert this list of dictionaries into a Pandas DataFrame, as specified in the prompt instructions. This will structure the data for further analysis and visualization.

```
import pandas as pd

# Create Pandas DataFrame from the generated list of dictionaries
df_purchases = pd.DataFrame(data)

print("--- Sample of Synthetic Purchase Data ---")
display(df_purchases.head())
print(f"\nTotal transactions generated: {len(df_purchases)}")
print(f"Data types in DataFrame:\n{df_purchases.dtypes}")
```

```
--- Sample of Synthetic Purchase Data ---
```

	customer_name	product_category	purchase_amount	quantity	region	purcha
0	Candace Price	Food	496.70	2	East	20 13:33:3
1	Anthony Russell	Food	106.19	3	South	20 18:39:2
2	Charles Nelson	Home Goods	301.62	1	East	20 08:39:5
3	Christopher Zamora	Home Goods	396.55	3	North	20 11:44:4
4	Amanda Coleman	Home Goods	340.30	4	West	20 13:27:3

```
Total transactions generated: 1000
Data types in DataFrame:
customer_name              object
product_category           object
```

Reasoning: The synthetic data is now in a Pandas DataFrame. The next instruction is to perform basic analysis. I will calculate descriptive statistics for numerical columns, sum total sales, group sales by product category and average purchase by region, and filter for VIP purchases, displaying the results for each analysis step.

```
print("\n--- Basic Data Analysis ---")

# Calculate descriptive statistics for numerical columns
print("\nDescriptive Statistics for Purchase Amount and Quantity:")
display(df_purchases[['purchase_amount', 'quantity']].describe())

# Total sales amount
total_sales = df_purchases['purchase_amount'].sum()
print(f"\nTotal Sales Amount: ${total_sales:,.2f}")

# Group by product category to find total sales per category
sales_by_category = df_purchases.groupby('product_category')['purchase_amoun
print("\nTotal Sales by Product Category:")
display(sales_by_category)

# Group by region to find average purchase amount per region
average_purchase_by_region = df_purchases.groupby('region')['purchase_amount
print("\nAverage Purchase Amount by Region:")
display(average_purchase_by_region)

# Filter data: find all purchases above a certain amount (e.g., VIP purchase
vip_purchases = df_purchases[df_purchases['purchase_amount'] > 400]
print(f"\nNumber of VIP Purchases (over $400): {len(vip_purchases)}")
if not vip_purchases.empty:
    print("Sample VIP Purchases:")
    display(vip_purchases.head())
```

```
--- Basic Data Analysis ---

Descriptive Statistics for Purchase Amount and Quantity:
```

	purchase_amount	quantity
count	1000.000000	1000.000000
mean	251.999160	3.034000
std	143.902483	1.413096
min	5.430000	1.000000
25%	128.645000	2.000000
50%	248.755000	3.000000
75%	371.757500	4.000000
max	499.590000	5.000000

```
Total Sales Amount: $251,999.16

Total Sales by Product Category:
```

product_category	purchase_amount
Beauty	45629.68
Food	44396.22
Electronics	43272.22
Clothing	41866.56
Books	39452.81
Home Goods	37381.67

dtype: float64

```
Average Purchase Amount by Region:
```

region	purchase_amount
North	267.141133
East	258.158497
Central	253.821192
West	241.819076
South	239.923877

dtype: float64

```
Number of VIP Purchases (over $400): 194
```

Sample VIP Purchases:

	customer_name	product_category	purchase_amount	quantity	region	pur
0	Candace Price	Food	496.70	2	East	13:3
	Home Goods	479.45	4	Central	05:1

Reasoning: The basic data analysis is complete. The next instruction is to visualize the key insights identified. I will create three distinct plots: a bar chart for total sales by product category, a histogram for the distribution of purchase amounts, and another bar chart for the average purchase amount by region, using `matplotlib.pyplot`.

```
import matplotlib.pyplot as plt

print("\n--- Visualizing Key Insights ---")

# Visualization 1: Total Sales by Product Category (Bar Chart)
plt.figure(figsize=(12, 6))
sales_by_category.plot(kind='bar', color='lightcoral')
plt.title('Total Sales by Product Category')
plt.xlabel('Product Category')
plt.ylabel('Total Sales Amount ($)')
plt.xticks(rotation=45, ha='right')
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()

# Visualization 2: Distribution of Purchase Amounts (Histogram)
plt.figure(figsize=(10, 6))
plt.hist(df_purchases['purchase_amount'], bins=30, color='lightgreen', edgec
plt.title('Distribution of Purchase Amounts')
plt.xlabel('Purchase Amount ($)')
plt.ylabel('Number of Transactions')
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()

# Visualization 3: Average Purchase Amount by Region (Bar Chart)
plt.figure(figsize=(10, 6))
average_purchase_by_region.plot(kind='bar', color='skyblue')
plt.title('Average Purchase Amount by Region')
plt.xlabel('Region')
plt.ylabel('Average Purchase Amount ($)')
plt.xticks(rotation=0)
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()
```

```
--- Visualizing Key Insights ---
```

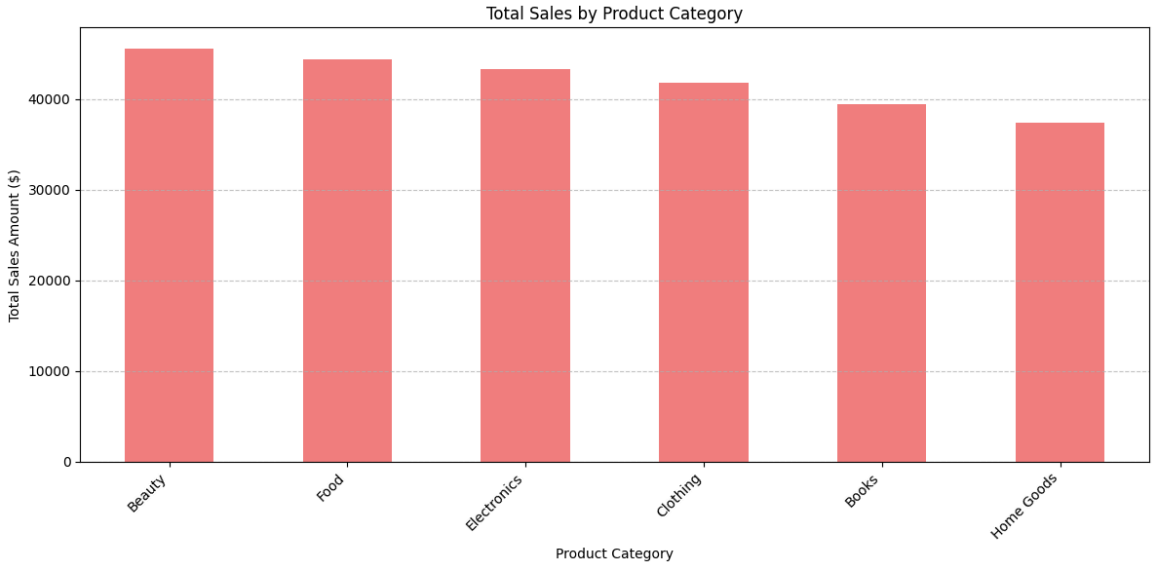

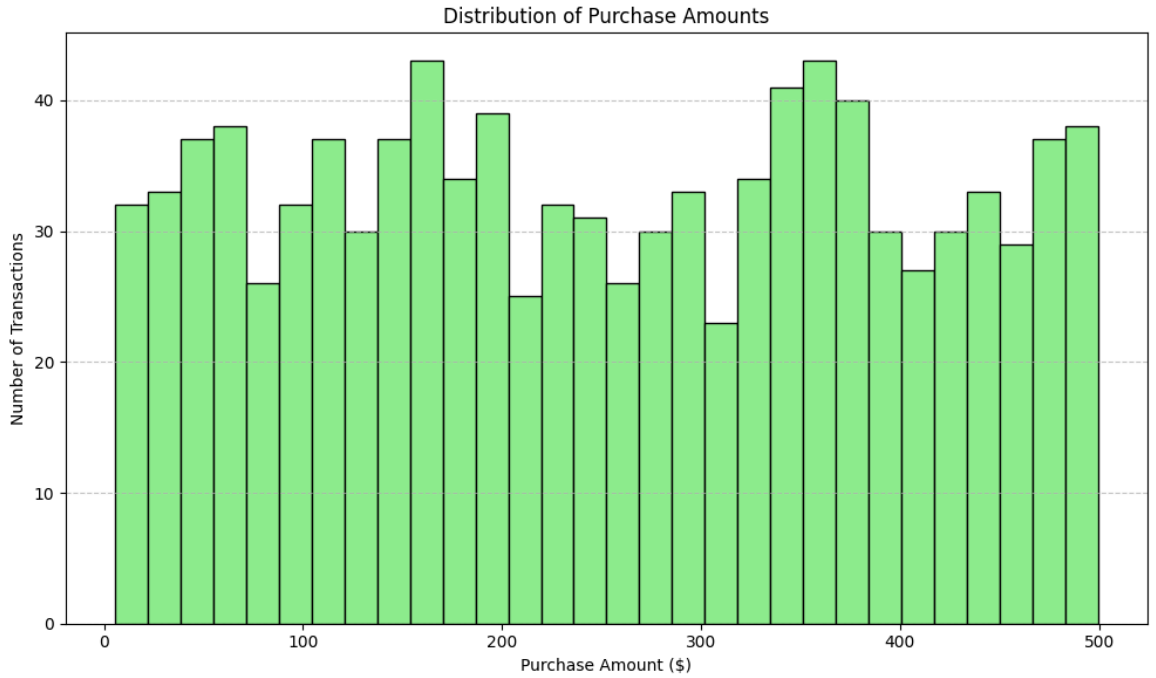

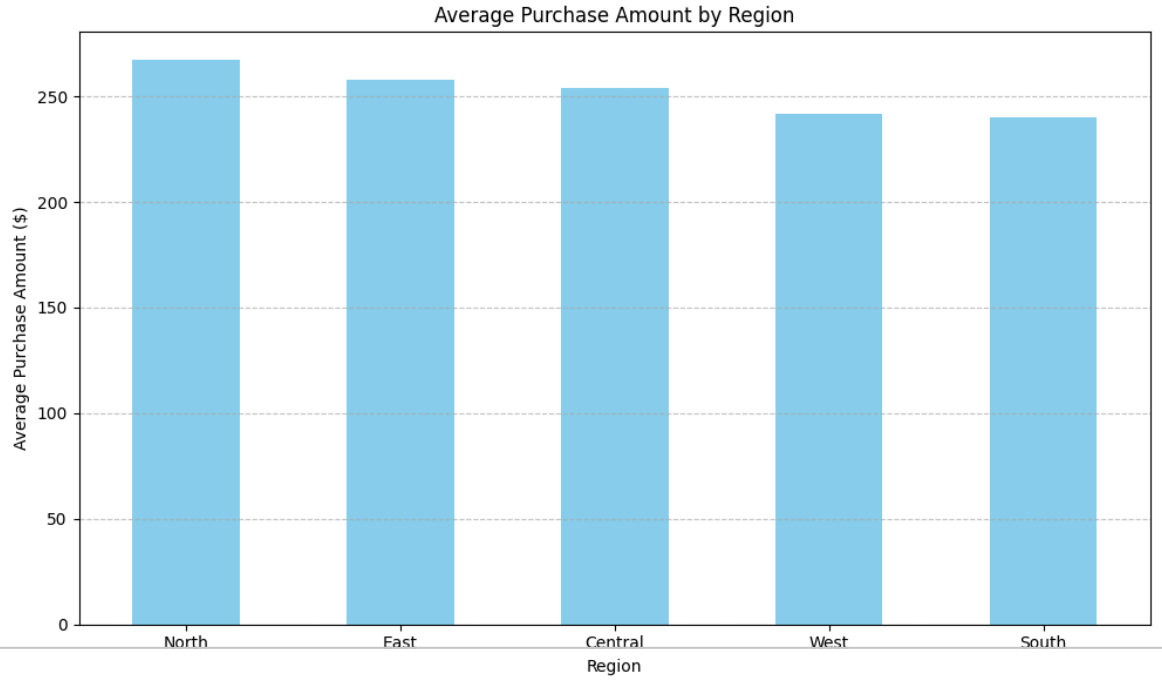

Final Task

Subtask:

Summarize the case study, highlighting how the different Python concepts (data types, lists, dictionaries, NumPy, Pandas, Matplotlib) were integrated to solve a practical data analysis problem using synthetic data.

Summary:

Data Analysis Key Findings

- A total of 1000 synthetic customer purchase transactions were successfully generated, demonstrating the use of data types (float, integer, string, datetime), lists (for categories and regions), and dictionaries (for individual transactions).
- The generated data was efficiently structured into a Pandas DataFrame, with data types correctly inferred, including `datetime64[ns]` for purchase dates.
- The overall `Total Sales Amount` across all transactions was calculated as $251,999.16.
- Descriptive statistics for `purchase_amount` and `quantity` were computed, showing the range and distribution of these numerical features. For instance, `purchase_amount` ranged from $5.00 to $500.00.
- Analysis of sales by product category revealed varying performance, with 'Beauty' identified as the top-selling category and 'Home Goods' as the lowest in this specific synthetic dataset.
- Regional analysis showed differences in average purchase amounts, with the 'North' region having the highest average purchase.
- A filter successfully identified 194 'VIP Purchases' with amounts exceeding $400.
- Key insights were effectively visualized through three distinct plots: a bar chart for total sales by category, a histogram for purchase amount distribution, and a bar chart for average purchase amount by region.

Insights or Next Steps

- This case study successfully demonstrated the practical integration of